


FUNDAMENTALS OF DEEP LEARNING

Regularization for Neural Networks

Did I practice the problems

or did I just **memorize** the solved examples ?

Take a deep breath, you're overfitting.



Shivang Kainthola

Regularization in Neural Networks

→ When a neural network or machine learning model performs too well on the training data but fails to **generalize to the testing data**, the problem is called **overfitting**.

→ Overfitting is a common issue while training deep neural networks or any machine learning models.

Overfitting in a neural network also manifests itself in or with other problems like :

- 1) Internal Covariate Shift
- 2) Co-Adaptation
- 3) Large Weights

→ To counter overfitting in neural networks, we use some **regularization techniques** :

- 1) Batch Normalization
- 2) Dropout
- 3) L1 Regularization
- 4) L2 Regularization

1) BATCH NORMALIZATION

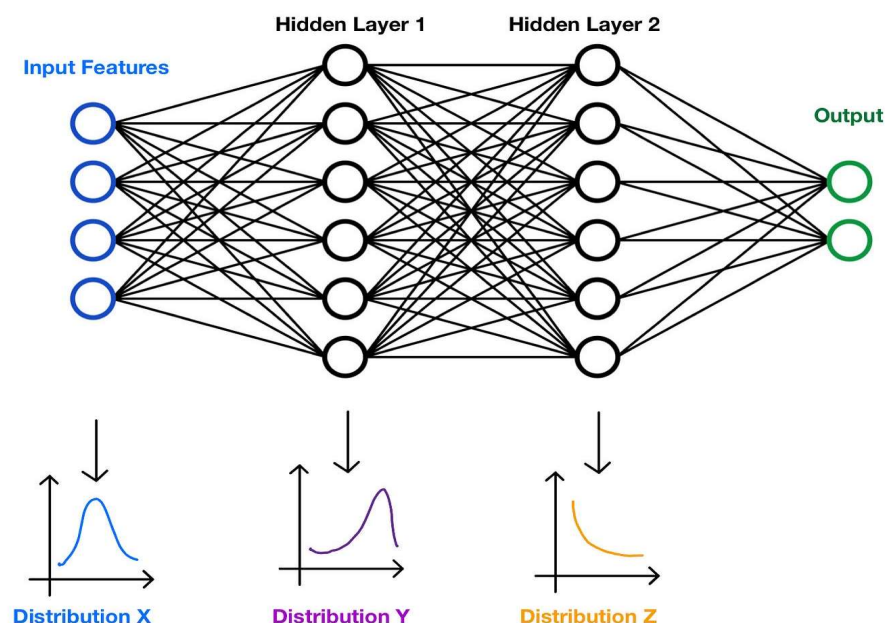
Problem : Internal Covariate Shift

During the training of a neural network by backpropagation, the parameters (weights and biases) are updated based on the error calculated by the loss function.

→ The activations of each layer depend on the inputs and parameters, both of which are changing during training.

→ Since the parameters as well as the inputs to each layer are constantly being updated, so there is **a shift in the distribution of inputs to each layer** which is called internal covariate shift.

Internal covariate shift **can slow down training and lead to instability.**



Solution : Batch Normalization


→ A **batch** or mini-batch is a collection of samples that will be passed through the network at one time for the weights update.

→ Batch Normalization is a **regularization technique** - where we normalize the inputs to a layer for every mini-batch.

→ Normalizing the data involves transforming it to have **mean** = 0 and **standard deviation** = 1.

→ Besides tackling internal covariate shift, it makes the gradient descent better.

<https://medium.com/@abheerchrome/batch-normalization-explained-1e78f7eb1e8a>



→ In a convolutional neural network, Batch Normalization is carried out with a **Normalization Layer** placed after the convolution layer.

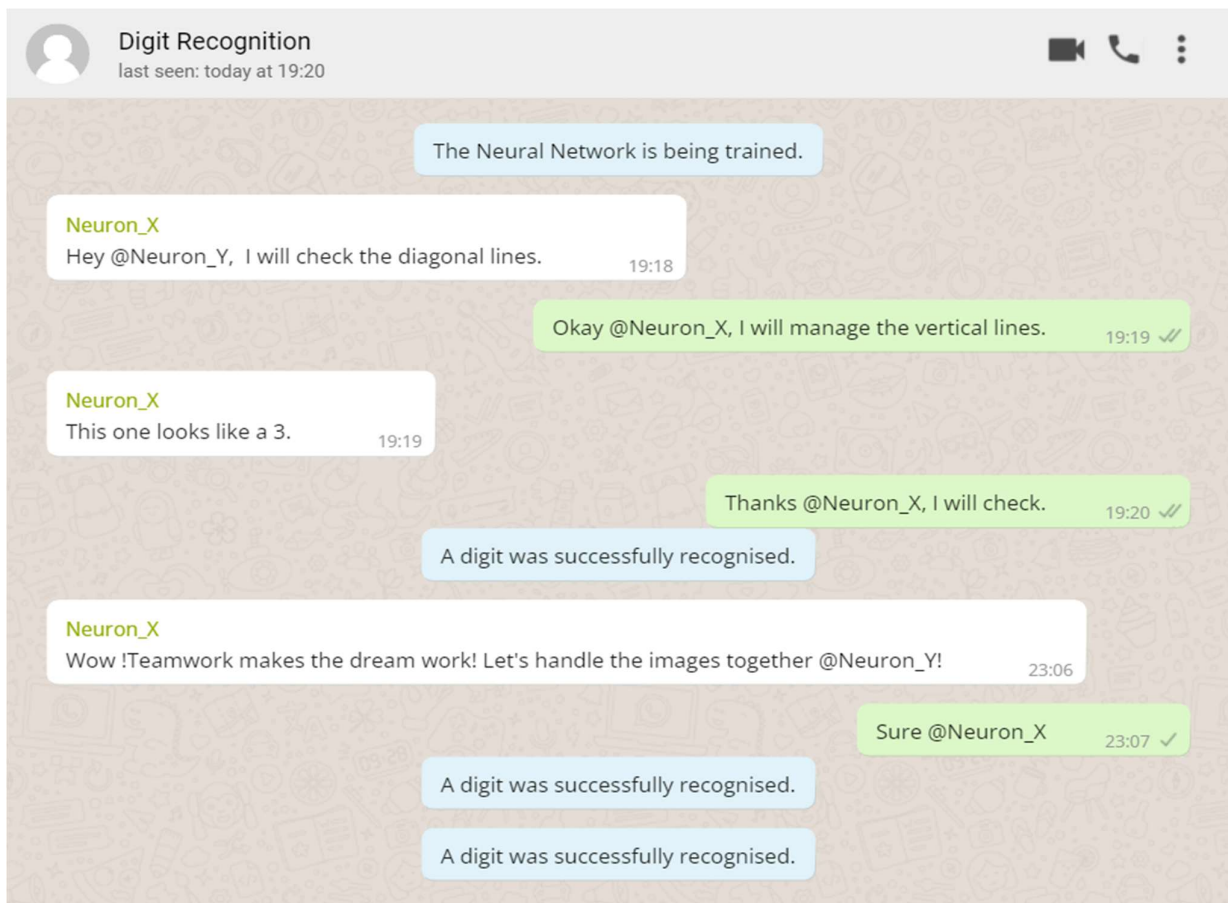
2) Dropout

Problem : Co-adaptation relationships

→ Co-adaptation is a situation when neurons become excessively reliant on one another.

→ To work properly, the neurons rely on the input of other co-adapted neurons.

→ With this co-adaptation relationship, the neurons also tend to be more generalized to training data, may underperform on testing data.



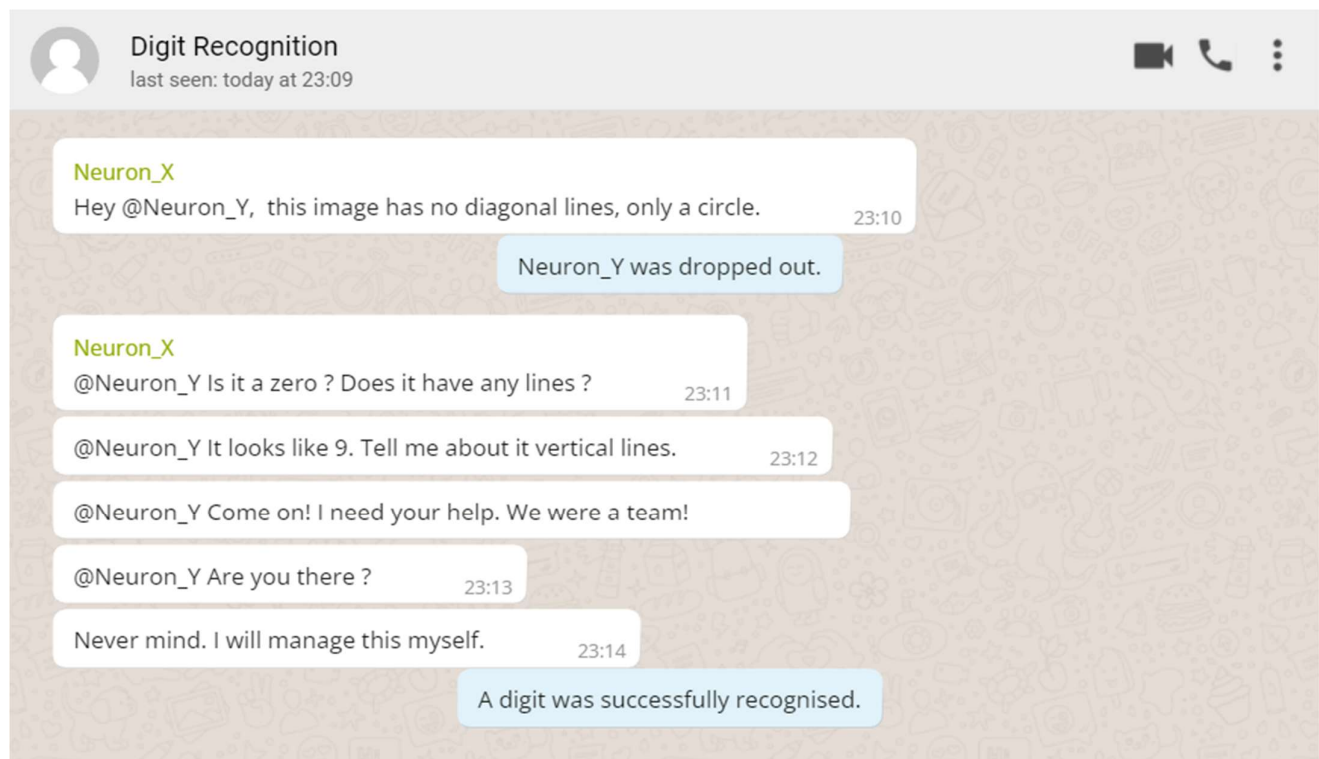
The co-adapted neurons X and Y can become dependent on each other.

Solution : Dropout

→ The ‘dropout’ method works **by randomly setting a fraction of the input units (neurons) in a layer to zero**, i.e. dropping them out, during each training iteration.

→ This dropped out fraction of neurons does not take part in forward pass (activation and gradient computation) or backpropagation (gradient updates) during training.

→ With dropout, the neurons are denied the convenience of making co-adaptations and relying on other neurons, since they must learn more robust features on their own.



→ Dropout method is applied to a neural network as a **Dropout()** layer, which takes the **fraction of neurons to be dropped out** as input.

3) L2 Regularization

Problem : Large weights

→ L2 Regularization is a popular regularization technique which penalizes models with large parameter (weight) values.

→ It works by **adding L2 regularization term to the loss function**, and when the loss function is minimized (by gradient descent - which seeks the global minima), it **steers the network away from having large weights**.

→ The loss function L of a neural network with L2 regularization term will be :

$\|w\|^2$ represents the squared L2 norm of the weights (sum of squares of all elements in the weight vector)

$$\mathbf{L}_{\text{L2 REGULARIZED}}(y, \hat{y}) = \mathbf{L}(y, \hat{y}) + \frac{1}{2} \lambda \|w\|^2$$

→ It is controlled by the parameter **lambda λ** , and regularization penalties are applied on a per layer basis.

4) L1 Regularization

→ L1 regularization penalizes large weights by adding the L1 regularization term to the loss function, similar in working to L2 regularization.

→ The loss function L of a neural network with L1 regularization term will be :

$\|w\|_1$ represents the L1 norm of the weights, which is the sum of the absolute values of all elements in the weight vector (w)

$$\mathbf{L}_{\text{L1 REGULARIZED}}(y, \hat{y}) = \mathbf{L}(y, \hat{y}) + \lambda \|w\|$$

→ It can be combined with L2 regularization, often known as Elastic Net regularization.

Thank You !

If you found this article helpful, please do leave a like and comment any feedback you feel I could use.

While I'm working on the next project, I'll be over at :



www.linkedin.com/in/shivang-kainthola-2835151b9/



shivang.kainthola64@gmail.com



<https://shivangkainthola28.medium.com/>



<https://github.com/HeadHunter28>

